

Computer Networks Assignment 3

NAVNEEL SINGHAL

November 21, 2020

Contents

1	Code	1
2	Part 1	7
3	Part 2	7
4	Part 3	8
5	Part 4	8
6	Graphs for <code>vayu.iitd.ac.in</code>	9
7	Graphs for <code>norvig.com</code>	10
8	Graphs for both servers at once	12
9	Graph for broken connections	16

1 Code

```
1  import socket
2  import hashlib
3  import collections
4  import threading
5  import sys
6  import csv
7  import time
8  import matplotlib.pyplot as plt
9
10  '''
11  HTTP/1.1 206 Partial Content
12  Date: Wed, 18 Nov 2020 06:30:10 GMT
13  Server: Apache/2.4.29 (Ubuntu)
14  Last-Modified: Mon, 22 Apr 2019 17:44:41 GMT
15  ETag: "63025a-5872205e3b440"
16  Accept-Ranges: bytes
17  Content-Length: 100
18  Vary: Accept-Encoding
19  Content-Range: bytes 0-99/6488666
20  Keep-Alive: timeout=5, max=100
21  Connection: Keep-Alive
22  Content-Type: text/plain
23
24  '''
25
26  def parseCompletely(clientSocket):
27      messageSize = int(1e9)
28      headerSoFar = bytearray(b'')
29      messageSoFar = bytearray(b'')
30      state = 0 # 0 if inside header and 1 if in body
31      while len(messageSoFar) < messageSize:
32          receivedMessage = clientSocket.recv(1024)
33          if receivedMessage == b'':
34              raise RuntimeError('Broken connection')
35          if state == 0:
36              headerSoFar += bytearray(receivedMessage)
```

```

37         pos = headerSoFar.find(b'\r\n\r\n')
38         if pos != -1:
39             messageSoFar = headerSoFar[pos + 4:]
40             headerSoFar = headerSoFar[:pos + 4].decode('ASCII')
41             lengthKeyPosition = headerSoFar.find('Content-Length: ')
42             assert lengthKeyPosition != -1
43             startValuePosition = lengthKeyPosition + len('Content-Length: ')
44             endValuePosition = startValuePosition
45             while headerSoFar[endValuePosition] != '\r':
46                 endValuePosition += 1
47             messageSize = int(headerSoFar[startValuePosition : endValuePosition])
48             state = 1
49             headerSoFar = bytearray(b'')
50         else:
51             messageSoFar += bytearray(receivedMessage)
52     return messageSoFar
53
54 def downloadAtOnce():
55     serverName = 'vayu.iitd.ac.in'
56     serverPort = 80
57     message = "GET /big.txt HTTP/1.1\r\nHost: vayu.iitd.ac.in\r\nConnection: keep-alive\r\n\r\n".encode('ASCII')
58     clientSocket = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
59     clientSocket.connect((serverName, serverPort))
60     clientSocket.send(message)
61     receivedMessage = parseCompletely(clientSocket)
62     clientSocket.close()
63     hashValue = hashlib.md5(bytes(receivedMessage)).hexdigest()
64     assert hashValue == '70a4b9f4707d258f559f91615297a3ec'
65
66 def downloadChunkWise():
67     serverName = 'vayu.iitd.ac.in'
68     serverPort = 80
69     message = "GET /big.txt HTTP/1.1\r\nHost: vayu.iitd.ac.in\r\nConnection: keep-alive\r\nRange: bytes="
70     messageEnd = "\r\n\r\n"
71     chunkSize = 1024 * 16
72     fileSize = 6488666
73     clientSocket = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
74     clientSocket.connect((serverName, serverPort))
75     receivedMessage = bytearray(b'')
76     missed = 0
77     for currentPtr in range(0, fileSize, chunkSize):
78         sendMessage = message + str(currentPtr) + '-' + str(min(currentPtr + chunkSize, fileSize) - 1) + messageEnd
79         clientSocket.send((sendMessage).encode('ASCII'))
80         while True:
81             try:
82                 x = parseCompletely(clientSocket)
83                 receivedMessage += x
84                 break
85             except:
86                 missed += 1
87                 clientSocket.close()
88                 clientSocket = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
89                 clientSocket.connect((serverName, serverPort))
90                 sendMessage = message + str(currentPtr) + '-' + str(min(currentPtr + chunkSize, fileSize) - 1) + messageEnd
91                 clientSocket.send((sendMessage).encode('ASCII'))
92     clientSocket.close()
93     hashValue = hashlib.md5(bytes(receivedMessage)).hexdigest()
94     assert hashValue == '70a4b9f4707d258f559f91615297a3ec'
95
96 def downloadPipelined():
97     serverName = 'vayu.iitd.ac.in'
98     serverPort = 80
99     message = "GET /big.txt HTTP/1.1\r\nHost: vayu.iitd.ac.in\r\nConnection: keep-alive\r\nRange: bytes="
100    messageEnd = "\r\n\r\n"
101    fileSize = 6488666
102    chunkSize = 1024 * 16 # 16 KB
103    pipelineSize = 10
104    clientSocket = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
105    clientSocket.connect((serverName, serverPort))
106    receivedMessage = bytearray(b'')
107    chunkSet = collections.OrderedDict()
108    totalChunks = 0
109    for currentPtr in range(0, fileSize, chunkSize):
110        chunkSet[(currentPtr, min(currentPtr + chunkSize, fileSize) - 1)] = None
111        totalChunks += 1
112    d = dict() # mapping from chunk ranges to chunks

```

```

113 currentChunkRange = None # a pair of integers determining the chunk
114 messageSoFar = bytearray(b'')
115 headerSoFar = bytearray(b'')
116 state = 0 # 0 for inside header and 1 for inside message
117 messageSize = int(1e9) # total message size - whenever the state is 0, change it to 1e9
118 completeChunks = 0 # whenever we read a chunk completely, add it to the dictionary and increment this variable by 1
119 currentlyParsing = False
120 totalRightNow = 0
121 received = min(totalChunks, pipelineSize)
122 while completeChunks < totalChunks:
123     while not currentlyParsing:
124         try:
125             if received == min(totalChunks - completeChunks, pipelineSize):
126                 received = 0
127                 sent = 0
128                 for (begin, end) in chunkSet:
129                     sendMessage = message + str(begin) + '-' + str(end) + messageEnd
130                     clientSocket.send((sendMessage).encode('ASCII'))
131                     sent += 1
132                     if sent == pipelineSize:
133                         break
134                 receivedMessage = bytearray(clientSocket.recv(1024))
135                 if receivedMessage == bytearray(b''):
136                     raise RuntimeError('Broken connection')
137                 currentlyParsing = True
138             except:
139                 currentChunkRange = None
140                 messageSoFar = bytearray(b'')
141                 headerSoFar = bytearray(b'')
142                 state = 0
143                 messageSize = int(1e9)
144                 clientSocket = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
145                 clientSocket.connect((serverName, serverPort))
146                 sent = 0
147                 for (begin, end) in chunkSet:
148                     sendMessage = message + str(begin) + '-' + str(end) + messageEnd
149                     clientSocket.send((sendMessage).encode('ASCII'))
150                     sent += 1
151                     if sent == pipelineSize:
152                         break
153             if state == 0:
154                 assert len(receivedMessage) != 0
155                 headerSoFar += receivedMessage
156                 pos = headerSoFar.find(b'\r\n\r\n')
157                 if pos != -1:
158                     receivedMessage = headerSoFar[pos + 4:]
159                     headerSoFar = headerSoFar[:pos + 4].decode('ASCII')
160                     lengthKeyPosition = headerSoFar.find('Content-Length: ')
161                     assert lengthKeyPosition != -1
162                     startValuePosition = lengthKeyPosition + len('Content-Length: ')
163                     endValuePosition = startValuePosition
164                     while headerSoFar[endValuePosition] != '\r':
165                         endValuePosition += 1
166                     messageSize = int(headerSoFar[startValuePosition : endValuePosition])
167                     rangeKeyPosition = headerSoFar.find('Content-Range: bytes ')
168                     assert rangeKeyPosition != -1
169                     startValuePosition = rangeKeyPosition + len('Content-Range: bytes ')
170                     endValuePosition = startValuePosition
171                     while headerSoFar[endValuePosition] != '/':
172                         endValuePosition += 1
173                     currentChunkRange = tuple([int(x) for x in headerSoFar[startValuePosition : endValuePosition].split('-')])
174                     state = 1
175                     headerSoFar = bytearray(b'')
176                     messageSoFar = bytearray(b'')
177                 else:
178                     receivedMessage = bytearray()
179                     currentlyParsing = False
180                     continue
181             if state == 1:
182                 i = 0
183                 while len(messageSoFar) < messageSize and i < len(receivedMessage):
184                     messageSoFar += receivedMessage[i : i + 1]
185                     i += 1
186                 receivedMessage = receivedMessage[i:]
187                 if len(messageSoFar) == messageSize:
188                     d[currentChunkRange] = messageSoFar

```

```

189         chunkSet.pop(currentChunkRange)
190         currentChunkRange = None
191         messageSoFar = bytearray(b'')
192         headerSoFar = bytearray(b'')
193         state = 0
194         messageSize = int(1e9)
195         completeChunks += 1
196         received += 1
197         if len(receivedMessage) == 0:
198             currentlyParsing = False
199             continue
200     ans = bytearray(b'')
201     for key in sorted(d):
202         ans += d[key]
203     hashValue = hashlib.md5(bytes(ans)).hexdigest()
204     assert hashValue == '70a4b9f4707d258f559f91615297a3ec'
205
206     '''
207     variables to be synchronized - chunkSet, completeChunks, d (not necessary since d is an array)
208     '''
209
210     lockSet = threading.Lock()
211     lockComplete = threading.Lock()
212
213     chunkSet = collections.OrderedDict()
214     totalChunks = 0
215     completeChunks = 0
216     fileSize = -1 #6488666
217     chunkSize = 1024 * 16
218     d = []
219     timeStamps = [[] for i in range(10000)]
220     startTime = time.time()
221
222     def chunkAllocator(pipelineSize):
223         global chunkSet
224         localChunkSet = collections.OrderedDict()
225         lockSet.acquire()
226         for (begin, end) in chunkSet:
227             localChunkSet[(begin, end)] = None
228             if len(localChunkSet) == pipelineSize:
229                 break
230         for (begin, end) in localChunkSet:
231             chunkSet.pop((begin, end))
232         lockSet.release()
233         return localChunkSet
234
235     def downloadSingleThread(threadNumber, serverName, pipelineSize, fileName):
236         global totalChunks, completeChunks, d, fileSize, chunkSize
237         serverPort = 80
238         message = "GET " + fileName + " HTTP/1.1\r\nHost: " + serverName + "\r\nConnection: keep-alive\r\nRange: bytes="
239         messageEnd = "\r\n\r\n"
240         chunkSize = 1024 * 16 # 16 KB
241         while True:
242             try:
243                 clientSocket = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
244                 clientSocket.settimeout(2)
245                 clientSocket.connect((serverName, serverPort))
246                 clientSocket.settimeout(None)
247                 break
248             except:
249                 continue
250         sizeProcessed = 0
251         headerSize = 0
252         receivedMessage = bytearray(b'')
253         currentChunkRange = None # a pair of integers determining the chunk
254         messageSoFar = bytearray(b'')
255         headerSoFar = bytearray(b'')
256         state = 0 # 0 for inside header and 1 for inside message
257         messageSize = int(1e9) # total message size - whenever the state is 0, change it to 1e9
258         currentlyParsing = False
259         totalRightNow = 0
260         localChunkSet = collections.OrderedDict()
261         while True:
262             #print(threadNumber)
263             lockComplete.acquire()
264             breakCondition = not (completeChunks < totalChunks)

```

```

265     lockComplete.release()
266     if breakCondition:
267         return
268     while not currentlyParsing:
269         try:
270             lockComplete.acquire()
271             breakCondition = not (completeChunks < totalChunks)
272             lockComplete.release()
273             if breakCondition:
274                 return
275             if len(localChunkSet) == 0:
276                 localChunkSet = chunkAllocator(pipelineSize)
277                 length = len(localChunkSet)
278                 if len(localChunkSet) == 0:
279                     return
280                 for (begin, end) in localChunkSet:
281                     sendMessage = message + str(begin) + '-' + str(end) + messageEnd
282                     clientSocket.send((sendMessage).encode('ASCII'))
283             receivedMessage = bytearray(clientSocket.recv(1024))
284             if receivedMessage == bytearray(b''):
285                 raise RuntimeError('Broken connection')
286             currentlyParsing = True
287         except:
288             try:
289                 currentChunkRange = None
290                 messageSoFar = bytearray(b'')
291                 headerSoFar = bytearray(b'')
292                 state = 0
293                 messageSize = int(1e9)
294                 clientSocket = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
295                 clientSocket.settimeout(2)
296                 clientSocket.connect((serverName, serverPort))
297                 clientSocket.settimeout(None)
298                 for (begin, end) in localChunkSet:
299                     sendMessage = message + str(begin) + '-' + str(end) + messageEnd
300                     clientSocket.send((sendMessage).encode('ASCII'))
301                 continue
302             except:
303                 continue
304         if state == 0:
305             assert len(receivedMessage) != 0
306             headerSoFar += receivedMessage
307             pos = headerSoFar.find(b'\r\n\r\n')
308             if pos != -1:
309                 receivedMessage = headerSoFar[pos + 4:]
310                 headerSoFar = headerSoFar[:pos + 4].decode('ASCII')
311                 lengthKeyPosition = headerSoFar.find('Content-Length: ')
312                 assert lengthKeyPosition != -1
313                 startValuePosition = lengthKeyPosition + len('Content-Length: ')
314                 endValuePosition = startValuePosition
315                 while headerSoFar[endValuePosition] != '\r':
316                     endValuePosition += 1
317                 messageSize = int(headerSoFar[startValuePosition : endValuePosition])
318                 rangeKeyPosition = headerSoFar.find('Content-Range: bytes ')
319                 assert rangeKeyPosition != -1
320                 startValuePosition = rangeKeyPosition + len('Content-Range: bytes ')
321                 endValuePosition = startValuePosition
322                 while headerSoFar[endValuePosition] != '/':
323                     endValuePosition += 1
324                 currentChunkRange = tuple([int(x) for x in headerSoFar[startValuePosition : endValuePosition].split('-')])
325                 state = 1
326                 headerSoFar = bytearray(b'')
327                 messageSoFar = bytearray(b'')
328                 headerSize = pos + 4
329             else:
330                 receivedMessage = bytearray()
331                 currentlyParsing = False
332                 continue
333         if state == 1:
334             i = 0
335             while len(messageSoFar) < messageSize and i < len(receivedMessage):
336                 messageSoFar += receivedMessage[i : i + 1]
337                 i += 1
338             receivedMessage = receivedMessage[i:]
339             if len(messageSoFar) == messageSize:
340                 d[currentChunkRange[0] // chunkSize] = messageSoFar

```

```

341         localChunkSet.pop(currentChunkRange)
342         sizeProcessed += messageSize + headerSize
343         timeStamps[threadNumber].append((time.time() - startTime, sizeProcessed))
344         currentChunkRange = None
345         messageSoFar = bytearray(b'')
346         headerSoFar = bytearray(b'')
347         state = 0
348         messageSize = int(1e9)
349         lockComplete.acquire()
350         completeChunks += 1
351         #print('Completed chunks, Thread number:', completeChunks, threadNumber)
352         lockComplete.release()
353         if len(receivedMessage) == 0:
354             currentlyParsing = False
355             continue
356     return
357
358 def getContentLength(clientSocket):
359     global fileSize
360     headerSoFar = ''
361     while True:
362         receivedMessage = clientSocket.recv(1024)
363         if receivedMessage == b'':
364             raise RuntimeError('Broken connection')
365         headerSoFar += (receivedMessage).decode('ASCII')
366         if headerSoFar.lower().find('404 not found') != -1:
367             fileSize = -1
368             return
369         pos = headerSoFar.find('\r\n\r\n')
370         if pos != -1:
371             headerSoFar = headerSoFar[:pos + 4]
372             lengthKeyPosition = headerSoFar.find('Content-Length: ')
373             assert lengthKeyPosition != -1
374             startValuePosition = lengthKeyPosition + len('Content-Length: ')
375             endValuePosition = startValuePosition
376             while headerSoFar[endValuePosition] != '\r':
377                 endValuePosition += 1
378             fileSize = int(headerSoFar[startValuePosition : endValuePosition])
379             return
380     fileSize = -1
381     return
382
383 def downloadMultipleConnections(l):
384     global totalChunks, completeChunks, d, fileSize, chunkSize
385     assert len(l) >= 1
386     # find fileSize by sending request to the first server in l
387     for ((serverName, filePath), _) in l:
388         # we can test for whether the resource is at these servers by sending a head request
389         # but we will assume that it exists, raising an exception if it is not at the first one
390         serverPort = 80
391         message = "HEAD " + filePath + " HTTP/1.1\r\nHost: " + serverName + "\r\nConnection: keep-alive\r\n\r\n"
392         clientSocket = None
393         while True:
394             try:
395                 getContentLength(clientSocket)
396                 break
397             except:
398                 try:
399                     clientSocket = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
400                     clientSocket.connect((serverName, serverPort))
401                     clientSocket.send(message.encode('ASCII'))
402                     continue
403                 except:
404                     continue
405         break
406     if fileSize == -1:
407         raise RuntimeError('File not found at the server', serverName)
408     print('File size:', fileSize)
409     d = [None for i in range(0, fileSize, chunkSize)]
410     for currentPtr in range(0, fileSize, chunkSize):
411         chunkSet[(currentPtr, min(currentPtr + chunkSize, fileSize) - 1)] = None
412         totalChunks += 1
413     pipelineSize = 1
414     t = []
415     threadCount = 0
416     for ((url, filePath), numThreads) in l:

```

```

417         t += [threading.Thread(target=downloadSingleThread, args=(threadCount + i, url, pipelineSize, filePath)) for i in range
418             threadCount + numThreads]
419     for th in t:
420         th.start()
421     for th in t:
422         th.join()
423     ans = bytearray(b'')
424     for k in d:
425         ans += k
426     hashValue = hashlib.md5(bytes(ans)).hexdigest()
427     print('File hash:', hashValue)
428
429 def parseUrl(url):
430     # http:// or https://
431     if len(url) >= 7 and url[:7] == 'http://':
432         url = url[7:]
433     elif len(url) >= 8 and url[:8] == 'https://':
434         url = url[8:]
435     pos = url.find('/')
436     assert pos != -1
437     return (url[:pos], url[pos:])
438
439 def main():
440     #downloadAtOnce()
441     #downloadChunkWise()
442     #downloadPipelined()
443     #l = [('vayu.iitd.ac.in', '/big.txt'), 10]
444     #l = [('norvig.com', '/big.txt'), 10]
445     assert len(sys.argv) > 1
446     with open(sys.argv[1], newline='') as csvfile:
447         reader = csv.reader(csvfile, delimiter=',')
448         l = [(parseUrl(line[0]), int(line[1])) for line in list(reader)]
449         downloadMultipleConnections(l)
450         plotData = [list(zip(*timeStamps[i])) for i in range(sum([line[1] for line in l]))]
451         for i in range(l[0][1]):
452             plt.plot(plotData[i][0], plotData[i][1], color='red')
453         if len(l) > 1:
454             for i in range(l[0][1], l[0][1] + l[1][1]):
455                 plt.plot(plotData[i][0], plotData[i][1], color='blue')
456         if len(sys.argv) > 2:
457             plt.savefig(sys.argv[2])
458         else:
459             plt.show()
460
461 if __name__ == '__main__':
462     main()

```

2 Part 1

Problem: Use TCP sockets to download a file in one go, ensuring that the file is received correctly by verifying file hashes.

Solution: The solution can be seen as in the function `downloadAtOnce()` in the code. It is quite simple; it creates a socket, connects, sends a request, receives the response, parses the message and closes the connection, and verifies the file hash as well to ensure that the file is received correctly (we always do this check whenever we download a file).

3 Part 2

Problem: Use the HTTP header fields to download the file in chunks.

Solution: The solution for this can be seen as in the function `downloadChunkWise()` in the code. For this, we iterate over all chunks, and use the same parser as in the previous function to parse the headers one at a time.

Note that in this function, we have no pipelining whatsoever, however, the function `downloadPipelined()` gives a pipelined implementation of the function. We have another parameter now - the number of chunks we request for in a window. To do a pipelined implementation, we have three states - when we are not parsing

(i.e., when we are receiving), when we are in the header, and when we are in the body of a chunk. I have taken care of the fact that the received message can have a header starting in the middle of the received message, and that the message might have multiple headers if the buffer size is a bit more than the chunk size.

The parsing is slightly different from the previous function, where we were guaranteed that there is only one header and one body in whatever is in the buffer. Here we maintain the current chunk range, the message so far, the header so far, the message size, completed chunks and whether we are parsing or not.

4 Part 3

Problem: Make the downloader more efficient by using threading to open multiple connections and pipelining.

Solution: There are a few changes in the pipelined code in the previous part to make the pipelined code thread-safe; here I have implemented a chunk allocator which allocates a few chunks to a thread from a global pool of available chunks when it is called. Also I have used some locks to ensure that all operations are thread-safe. Some parameters are shifted to the thread arguments (for making it easier to call them from the function that creates all these threads). The rest of the code is pretty much the same and can be found in the function `downloadSingleThread()` for one thread, and the thread manager is the function `downloadMultipleConnections()`.

To implement the downloads from different servers, we first parse the URLs (which should contain the path of the file in the end) and then for the specified number of threads, use that information to download the chunks. To avoid hardcoding the file name and the URL, and hence the file size, we assume that the file actually exists at those paths and send a HEAD request to find the file size (we can check for file existence using a simple check which is done for the first server).

The graphs for some of the examined cases are attached towards the end.

1. **Does your download time keep decreasing with more and more parallel TCP connections?**

Yes, it decreases with more parallel TCP connections. This happens because the idle time due to blocking calls is now occupied by threads that are free, and thus making the download more efficient.

2. **Are some connections faster than the others? Do some connections get stalled for some reason?**

Within a certain amount of variance, all connections have roughly the same speeds and connections do not get stalled.

3. **Does your download time reduce if you download from different sources in parallel?**

Keeping the total number of threads constant, the situation where we have more threads to a faster server has a better download speed. However, adding more threads of another server almost always reduces the download time (if the number of threads of the original server are kept the same).

4. **What does this tell you about where the bottlenecks lie?**

In the case of the `vayu.iitd.ac.in` server, the connection is quite fast. Since python allows only one thread to run at a time, and there is not much effect of increasing the threads beyond a certain point, it seems that the RTT is not the bottleneck in the case of this server. However, for the `norvig.com` server, the RTT is the main bottleneck, since on increasing the number of threads, the performance goes up in a manner proportional to the number of threads, and the main reason where the threads reduce the time is when there is a blocking call, since python virtual machine runs a thread one at a time.

5. **Is one server faster than the other? Is your program able to use this to download more from the faster server**

Yes, the `vayu.iitd.ac.in` server is faster than the `norvig.com` server, as can be seen from ping response times (8ms and 200+ms). The program is able to use this and since the total time taken by all threads is the same, we see that it roughly divides the load in the ratio of the per-thread download speeds.

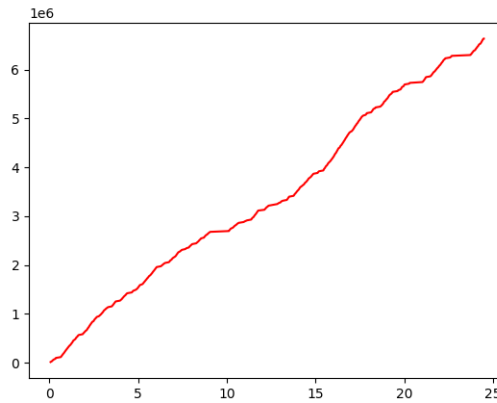
5 Part 4

Problem: Make this downloader more resilient by keeping a check for connection failures etc.

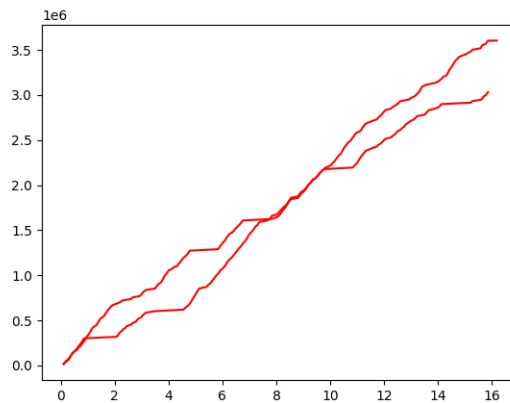
Solution: Now to make this resilient, we have to take care of disconnections, as well as errors in which the `recv` and `send` calls fail or connection is broken (in which case `recv` calls send an empty response). To ensure that these are taken care of, we enclose these calls in `try-except` blocks, and whenever they fail, we create a new socket and resend the remaining pipelined requests again.

6 Graphs for `vayu.iitd.ac.in`

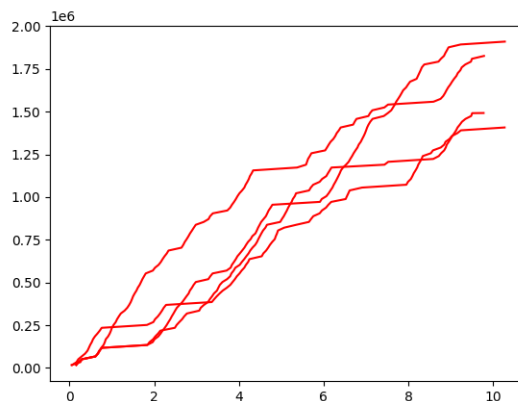
As can be seen below, the time decreases as we increase the number of TCP connections. The load shared by the threads is roughly equal, and the time taken by each thread to stop is also roughly the same.



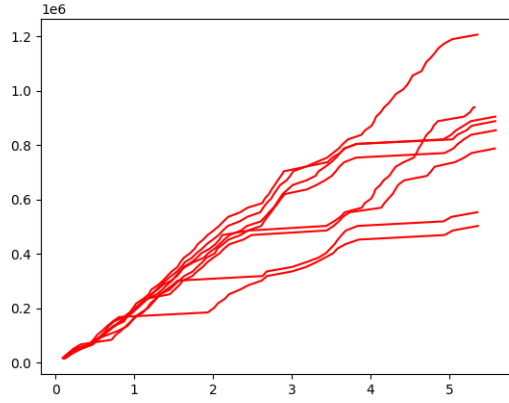
Bytes downloaded v/s time for 1 thread



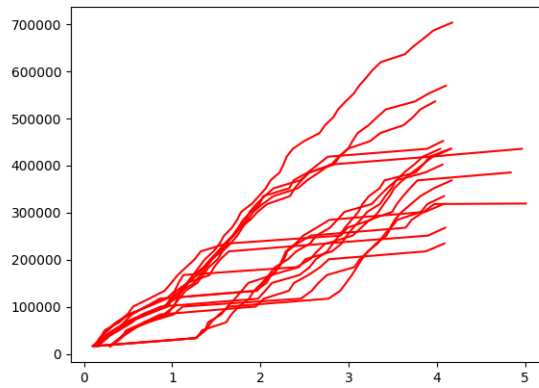
Bytes downloaded v/s time for 2 threads



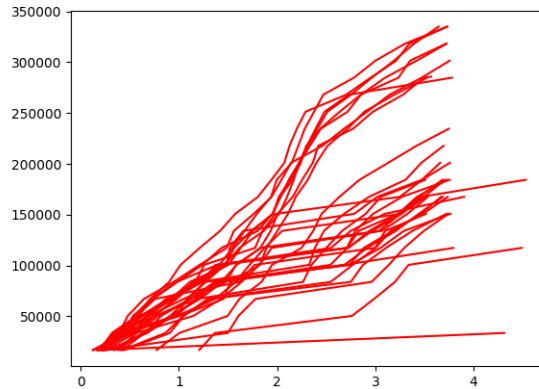
Bytes downloaded v/s time for 4 threads



Bytes downloaded v/s time for 8 threads



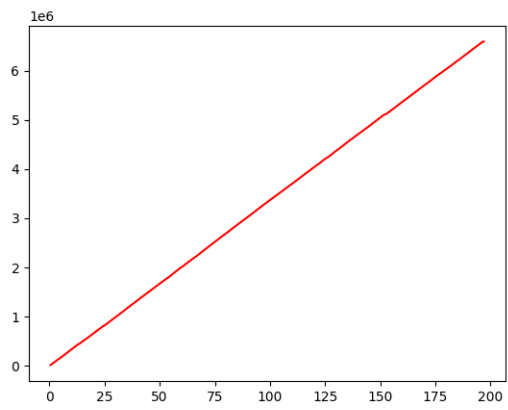
Bytes downloaded v/s time for 16 threads



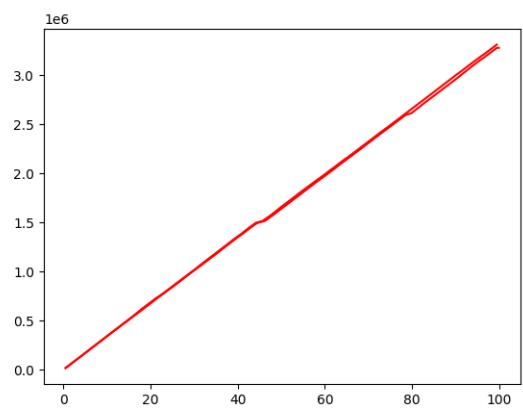
Bytes downloaded v/s time for 32 threads

7 Graphs for norvig.com

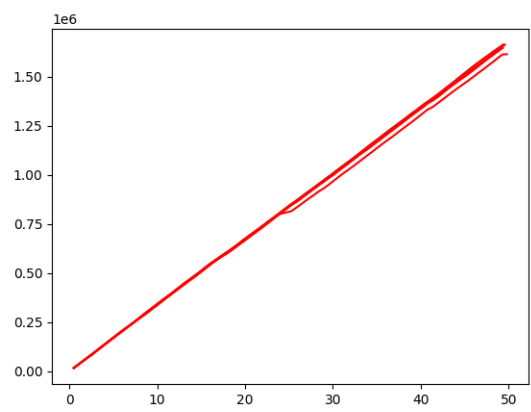
As for the previous server, here we observe a similar trend, but the time taken is almost inversely proportional to the number of threads. This happens since the RTT is the bottleneck here (as was seen on my connection, a ping to this server takes roughly 250ms, while that to the previous one takes 8ms). Most of the randomness is also gone, since the threads wait and finish waiting about the same time (there is less relative variance in the RTT for this server since the RTT is large, so threads are allocated more fairly).



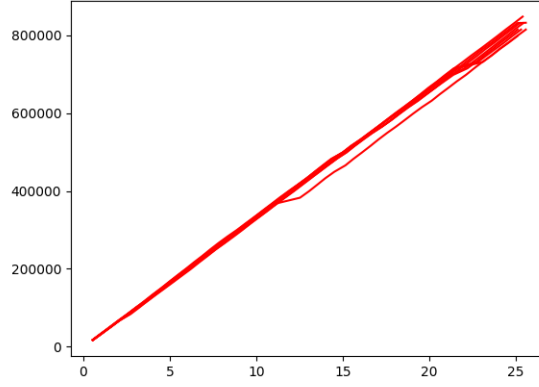
Bytes downloaded v/s time for 1 thread



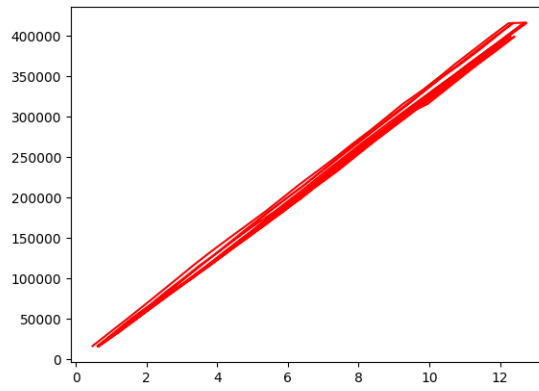
Bytes downloaded v/s time for 2 threads



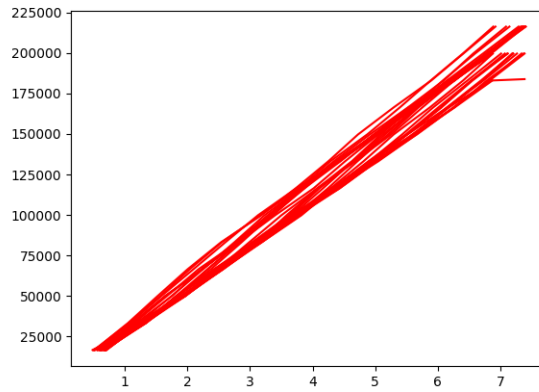
Bytes downloaded v/s time for 4 threads



Bytes downloaded v/s time for 8 threads



Bytes downloaded v/s time for 16 threads

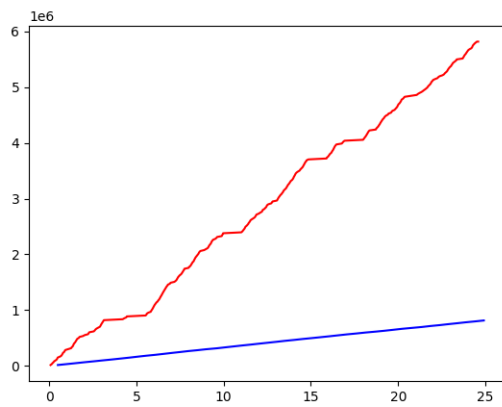


Bytes downloaded v/s time for 32 threads

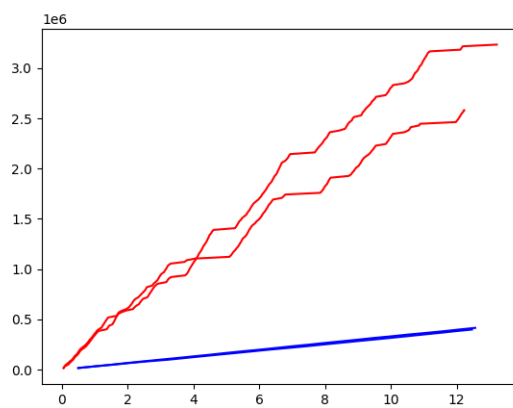
8 Graphs for both servers at once

From various runs while changing the number of threads, by keeping the number of threads same for both servers, and then keeping the total number of threads constant, we can notice that the total size downloaded by the threads of a server is proportional to the speed of the server with that number of threads when comparing across servers. This shows that the distribution among threads is nearly optimal with respect to total time taken, and all the threads take the same time.

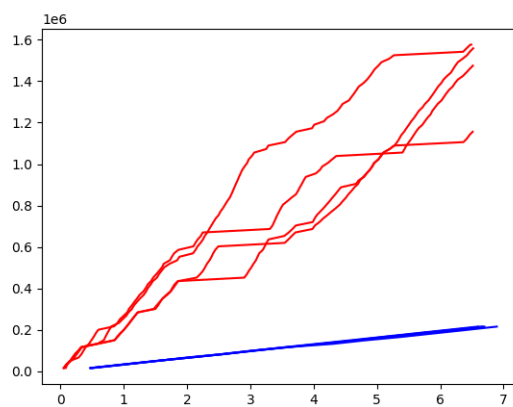
We also see that when we have the same total number of threads, the total download speed is higher when we have more threads for the faster server, as can be expected simply.



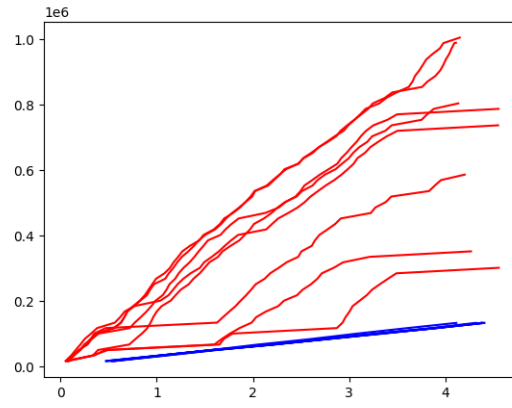
Bytes downloaded v/s time for 1 thread each



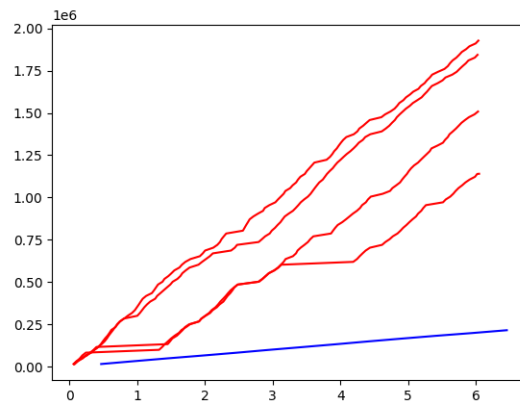
Bytes downloaded v/s time for 2 threads each



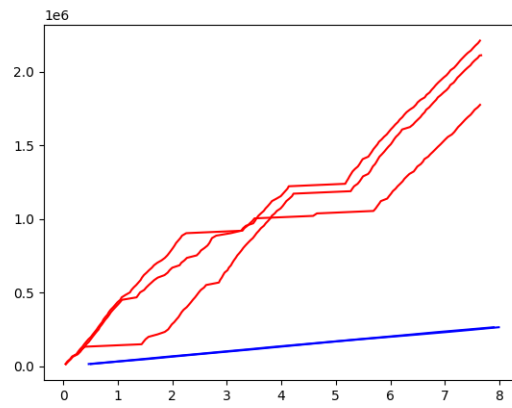
Bytes downloaded v/s time for 4 threads each



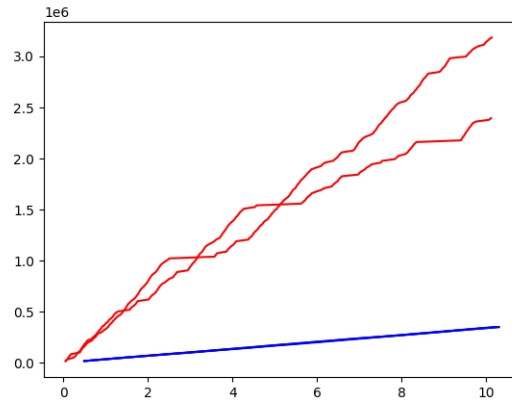
Bytes downloaded v/s time for 8 threads each



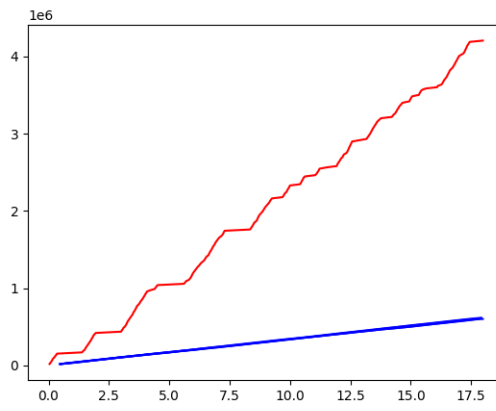
Bytes downloaded v/s time for 4, 1 threads on the first, second server respectively



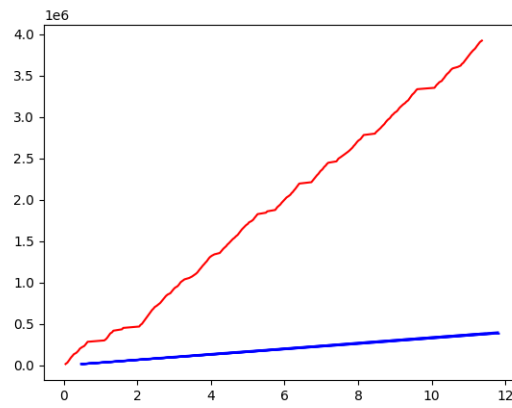
Bytes downloaded v/s time for 3, 2 threads on the first, second server respectively



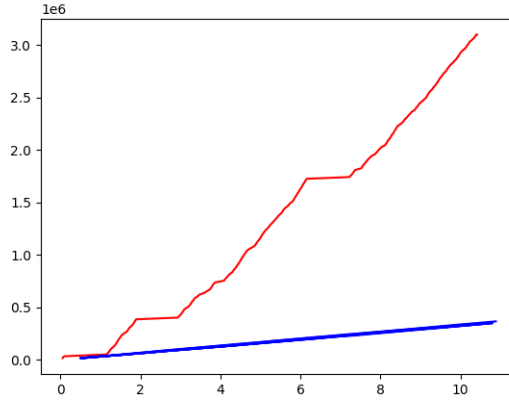
Bytes downloaded v/s time for 2, 3 threads on the first, second server respectively



Bytes downloaded v/s time for 1, 4 threads on the first, second server respectively



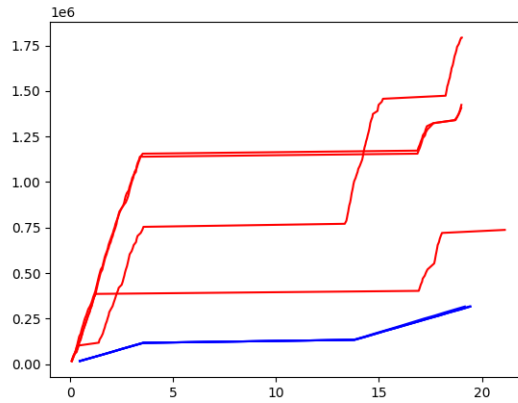
Bytes downloaded v/s time for 1, 7 threads on the first, second server respectively



Bytes downloaded v/s time for 1, 10 threads on the first, second server respectively

9 Graph for broken connections

As can be seen in the graph, broken connections do not fatally crash the download, and the download resumes when we come back online. Note that the alignment is not that clear, since we push the (timestamp, bytes downloaded) pairs only when we complete downloading a chunk, and we don't sample the downloaded bytes at a fixed interval of time. The connection timeout was set to be 2 seconds, and hence this also contributes to such an alignment.



Bytes downloaded v/s time for a connection broken due to disconnection from the internet